# Deep Neural Networks

## Activation Functions and Depth Analysis

Tigist W.

# Introduction

- This mini project explores the impact of different activation functions and the role of network depth in ANNs.
- We compare Sigmoid, Tanh, and ReLU activations.
- Investigate ReLU variants: Leaky ReLU and Parametric ReLU.
- Analyze network performance for different depths.

# Activation Functions Overview

- Sigmoid: Output range (0,1); suffers from vanishing gradients.
- Tanh: Output range (-1,1); still suffers from vanishing gradients.
- ReLU: Allows unbounded positive values; mitigates vanishing gradient issues.
- Leaky ReLU & Parametric ReLU: Fix the dying ReLU problem by allowing small negative slopes.

# Experiment Setup

- **Dataset: MNIST (Handwritten digits classification)**
- **Model: Fully Connected DNN**
- **Evaluation: Loss, gradient magnitude, and accuracy**
- **Tools: PyTorch, Matplotlib for visualizations**

# Activation Function Comparison

## Gradient Magnitude Analysis

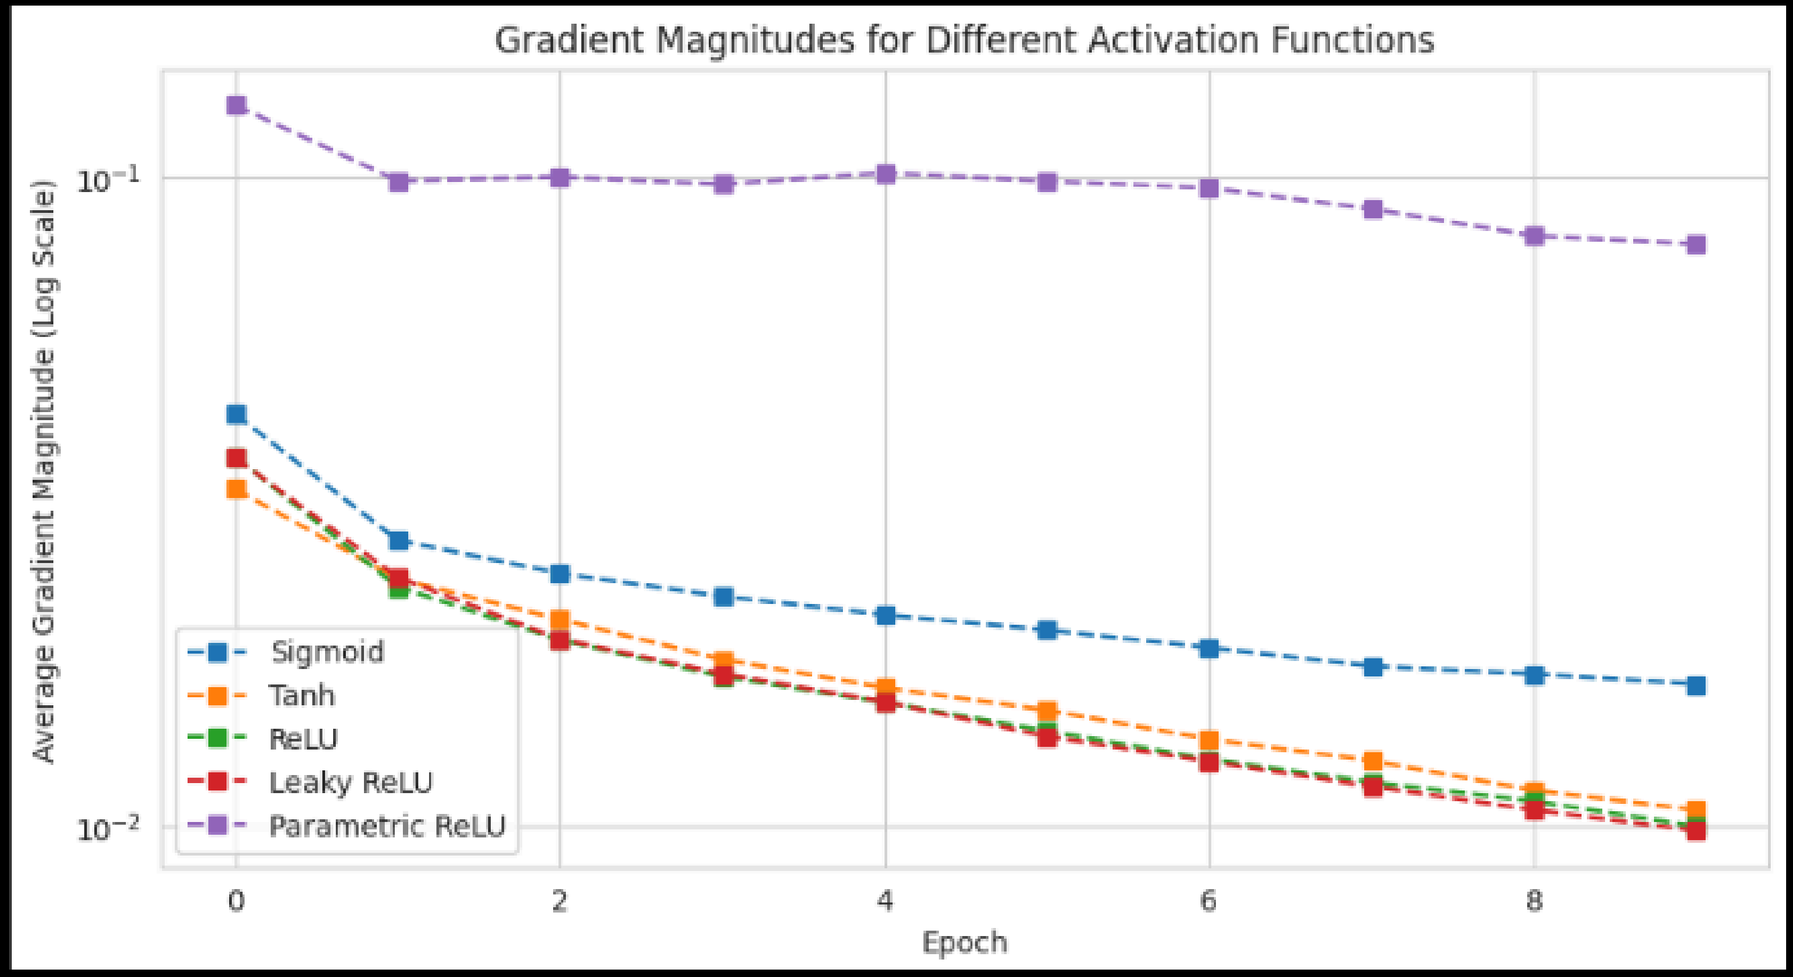

# Performance Visualization

**Loss curves: Show that ReLU converges faster than Sigmoid & Tanh.**

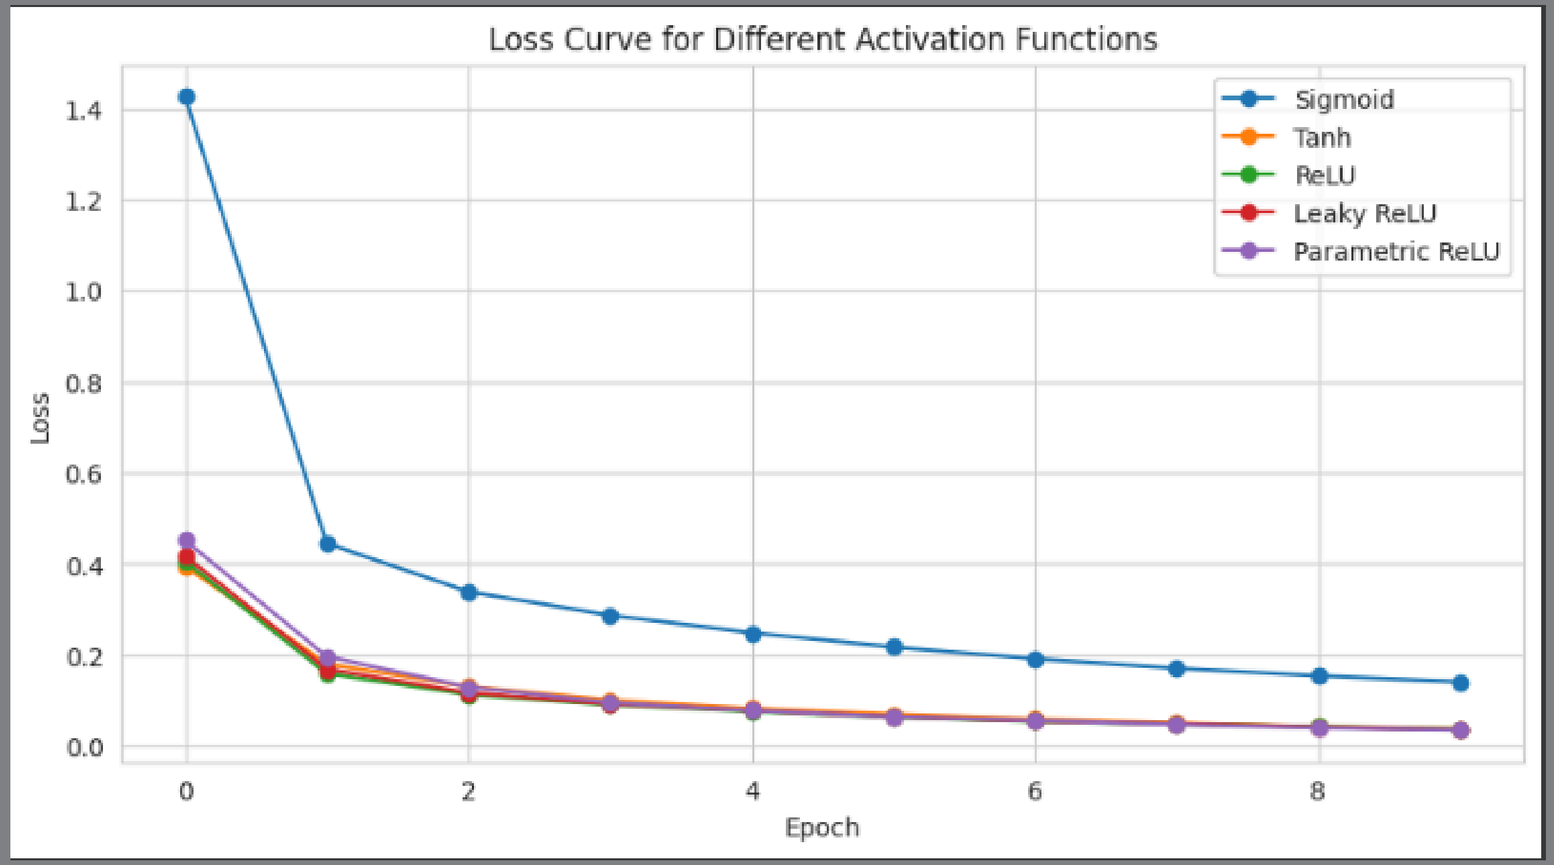

# Depth of ANNs

- **Testing different network depths:**
    - **2-layer shallow network**
    - **3-layer moderate network**
    - **5-layer deep network**
    - **7-layer very deep network**
    - **9-layer extra deep network**

# Depth vs. Performance

- Shallow networks: Struggle with feature extraction.
- Deep networks (5-7 layers): Achieve optimal accuracy.
- Extra deep networks (10+ layers): Show degradation due to vanishing gradients & overfitting.

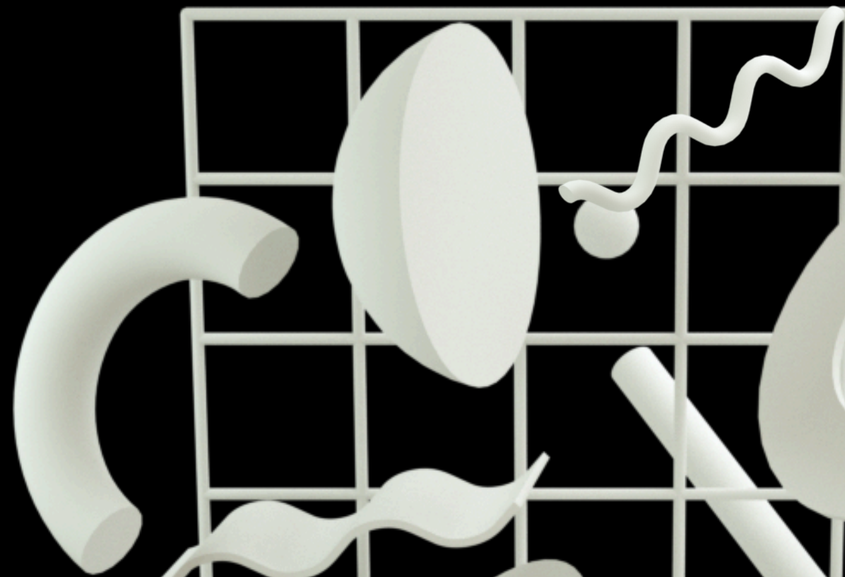

# Depth vs. Performance

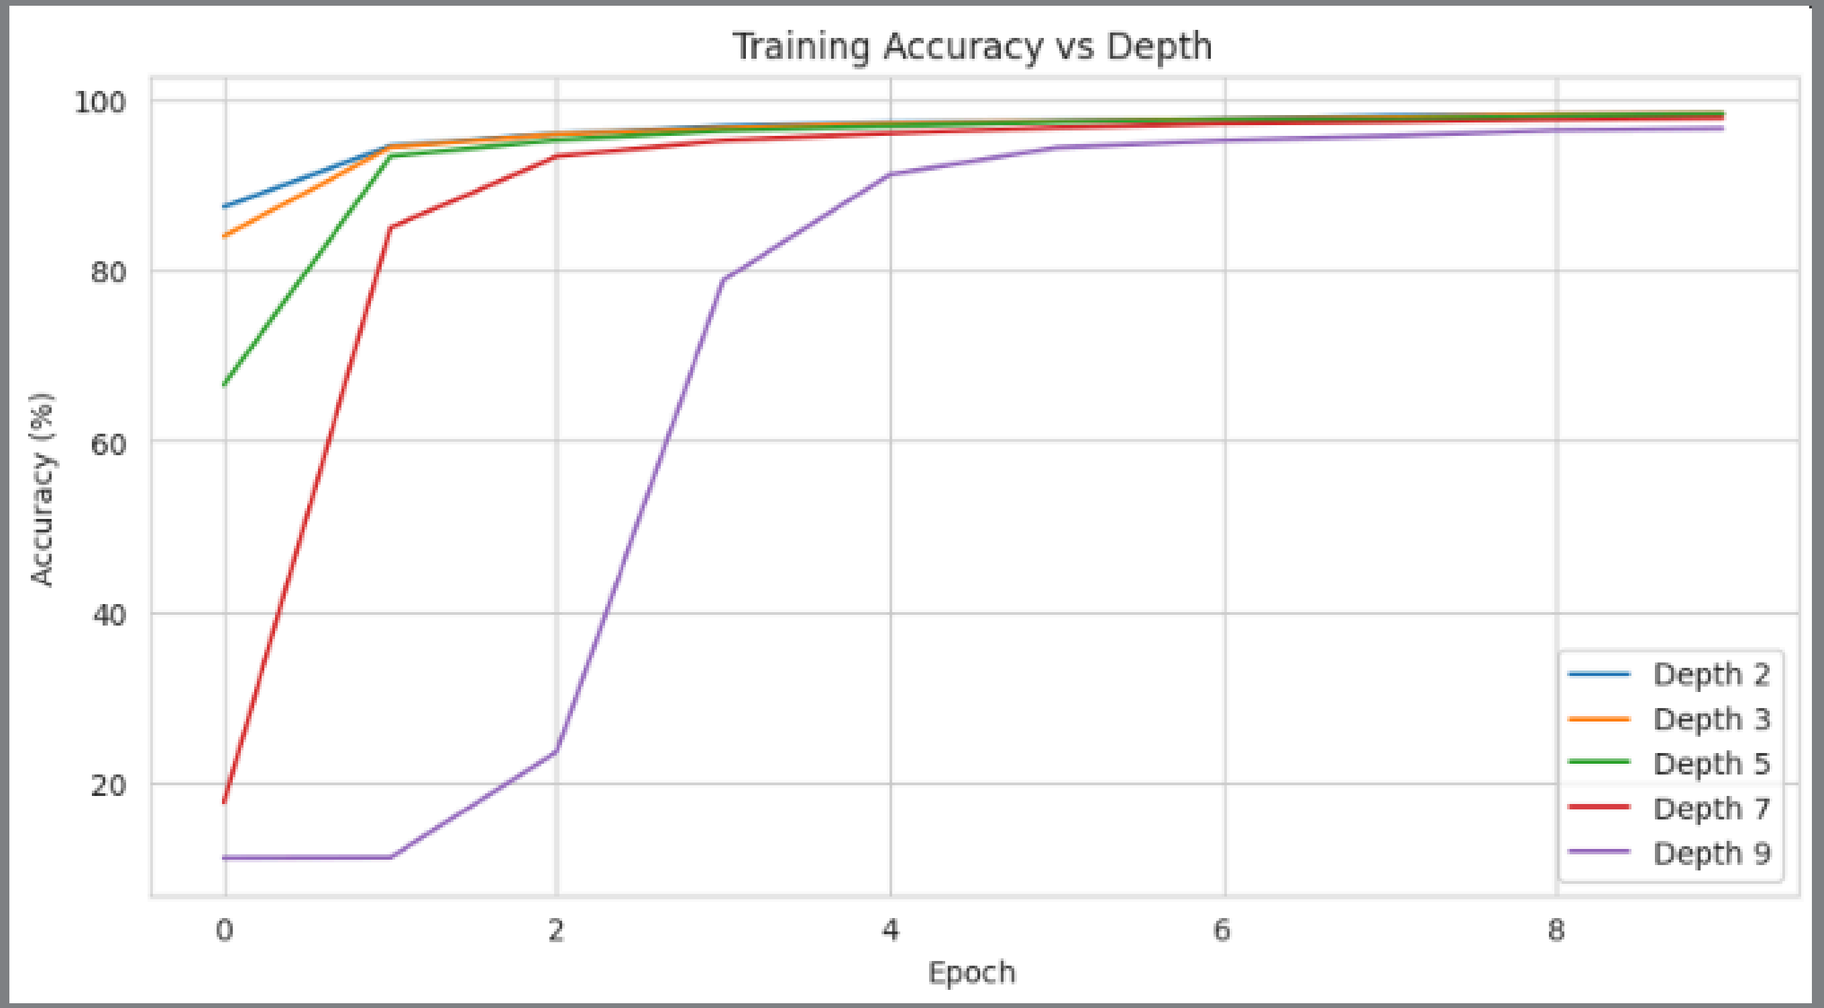

# Thank You!